

Snoop-Based Multiprocessor Design

Module-4



Hemangee K. Kapoor
Department of CSE
IIT Guwahati

Hemangee K. Kapoor

Single-Level caches with an Atomic Bus (1)

Module-4

Lecture-2



Hemangee K. Kapoor

List of design issues

- Here we have some assumptions but they are physically realistic
- List of preliminary design issues
 - 1) Design of cache controller and tags. Both processor and bus need to lookup
 - 2) How and when to present snoop results on the bus
 - 3) Dealing with write-backs, as they may cause race conditions
 - 4) Overall set of operations for memory access are not atomic. This can introduce race conditions
 - 5) Any issues that may arise wrt deadlock, livelock, starvation, serialisation, etc.

Hemangee K. Kapoor

(i) Cache controller + Tag design

- Uniprocessor cache
 - } Data, tags, state, comparator, controller and bus interface
- Processor cache interaction
 - } Compare tags, hit/miss, update state bits, may need block transfer to/from memory
- Cache controller initiates bus operation
 - } Assert request for bus; Wait for bus grant; Drive address and command lines; Wait for command to be accepted by relevant device; Transfer data

Hemangee K. Kapoor

Extend uniprocessor cache

- Cache controller is = FSM
- Sequence of steps in bus transactions = FSM
- Note that coherence protocol is another block-level FSM
- Changes
 - › Monitor bus
 - › Respond to processor request
- To implement snooping protocol, cache controller has bus-side and processor-side FSM


Hemangee K. Kapoor

Duplicate tags

- Processor-side compares tag
- Snoop also compares tags to take appropriate actions
- Cache controller sends Bus transactions in both cases
- Tag comparison needed by both
- Therefore dual-tags (don't duplicate data) or use dual-ported RAM for tags
 - › Minimal stall on tag updates. General operation of tag compare is fast

Hemangee K. Kapoor


Duplicate tags + FSMs



Hemangee K. Kapoor

Cache controller extends

- Cache controller also responds to bus transactions. Has to compare bus address with own address
- If hit = may have to put data on bus
- For update protocol = snoop data from bus and update own cache



Hemangee K. Kapoor

(ii) Reporting Snoop Results: When?

- All caches snoop on bus and do tag compare. Collective result of (tag compare) snoop must be sent on bus before transaction can proceed
- This is required as Memory must know if it should give data. i.e. Memory or cache will give
- Keep decision delay to minimum
- Three options:
 - › Fixed Delay
 - › Variable Delay
 - › Immediately

Hemangee K. Kapoor

Snoop results in Fixed Delay

- After fixed number of clock cycles after the address appears on bus
- This needs dual-tags to reduce contention with processor
- Still must be conservative, as CPU may lock both tags arrays when updating state (e.g. 'E' -> 'M')
- Advantage is main memory design not affected
 - › Cache-to-cache handshake is simple
- Dis-advantage is extra hardware, Potentially longer latency
- Ex: Pentium Pro, HP Server, Sun Enterprise

Hemangee K. Kapoor

Snoop results in Variable Delay

- Memory assumes cache will supply data, until all caches say otherwise
- Less conservative, as we need not assume worst case delay to compute snoop results
- More flexible, more complex (handshakes required between cache <--> memory)
- Memory can however start fetching data. If cache gives data then memory data is not used and memory access stopped
- Ex: SGI Challenge

Hemangee K. Kapoor

Snoop results Immediately

- Memory maintains 1-bit per block to indicate if it is being modified or not by any cache
- This bit helps memory to decide if memory should send data. Need not wait for snoop results
- Dis-advantage is extra hardware complexity to memory sub-system


Hemangee K. Kapoor

(iii) Reporting Snoop Results: How?

- Collective response from caches must appear on the bus
- Ex: in MESI protocol, we need to know
 - › Is block dirty?
 - › Should memory respond with data or some cache will provide data
 - › Is block shared?
 - › This is required to decide if one should load in 'E' or 'S' state

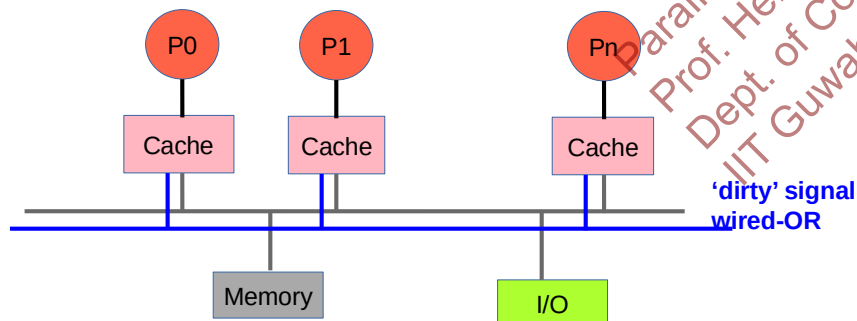
Hemangee K. Kapoor

Shared signal



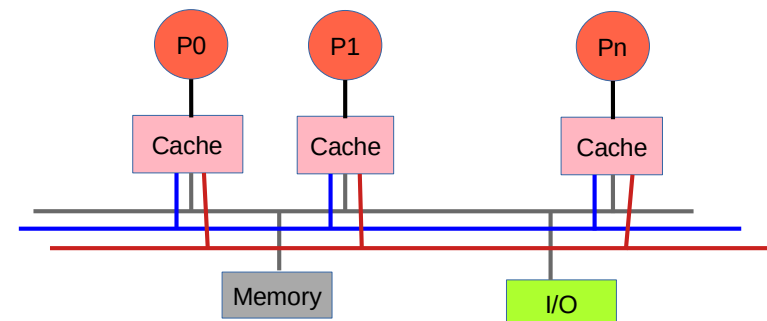
Hemangee K. Kapoor

Dirty signal



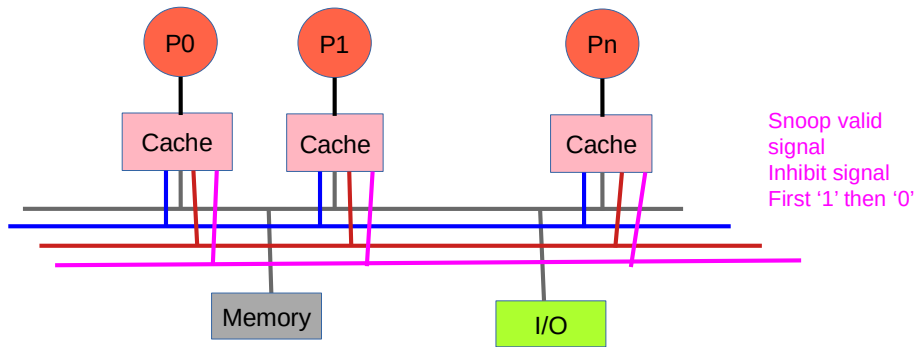
Hemangee K. Kapoor

Shared+ dirty signal



Hemangee K. Kapoor

Snoop-valid signal



Hemangee K. Kapoor

Wired-OR signals

- Use three wired-OR signals
 - Shared**: asserted if any cache (except Requestor) has a copy
 - Dirty**: asserted if any cache has modified copy. Need not know which, since it will know what to do
 - Snoop-Valid**: this is an **inhibit signal**. Asserted until all caches have completed their snoop. When de-asserted, Memory and Requestor can examine other two signals

Hemangee K. Kapoor

Who provides data?

- Full Illinois-MESI is more complex as
 - Block is provided by cache and not Memory
 - Then priority scheme as to which cache gives block if multiple have it
- Therefore commercial systems avoid cache-to-cache transfer
- Ex: SGI Challenge and Sun Enterprise use cache transfers only for modified data
 - SGI Challenge --> updates Memory when cache gives copy
 - Sun Enterprise --> uses 'O' bit. MOESI and Memory does not update when cache gives


Hemangee K. Kapoor

(iii) Dealing with write-backs

- At block replacement => 2 Bus transactions
 - Bring new block + write-back old block
- Want to reduce processor wait time on a cache miss
- Therefore service the miss first and then do write-back asynchronously
- Need additional storage : write-back buffer

Hemangee K. Kapoor


Write buffer



Hemangee K. Kapoor

(iii) Dealing with write-backs

- Bus transactions for this block can come. So snoop has to check even the write-back buffer
- Cache controller may have to cancel the write-back, if the block is used from the write-back buffer



Hemangee K. Kapoor

Thank you

Parallel Computer Architecture
 Prof. Hemangee K. Kapoor
 Dept. of Computer Science And Engineering
 IIT Guwahati

Hemangee K. Kapoor